

Kernel Methods

Kernel embedding of distributions and large scale kernel machines

Hachem Kadri, Liva Ralaivola
prenom.nom@lis-lab.fr

LIS, CNRS, Aix-Marseille Université, France



Outline

Kernel Embedding of Distributions

Example: the 2-sample test problem

Kernel Embedding of Distributions

Large Scale Kernel Methods

Random Features for Kernel Machines

Fast optimization methods

Outline

Kernel Embedding of Distributions

Example: the 2-sample test problem

Kernel Embedding of Distributions

Large Scale Kernel Methods

Random Features for Kernel Machines

Fast optimization methods

Two-Sample Problem (see [Gretton et al., 2007, Gretton et al., 2012])

Problem

- ▶ Let $S_X = \{x_i\}_{i=1}^{n_x} \sim P$ and $S_Y = \{y_i\}_{i=1}^{n_y} \sim Q$
- ▶ Question: given S_X and S_Y , is it possible to determine whether $P = Q$?
- ▶ Or, is it possible to provide a statistical test to decide $P = Q$:
 - $H_0 : P = Q$ (null hypothesis)
 - $H_A : P \neq Q$ (alternative hypothesis)

Two-Sample Problem (see [Gretton et al., 2007, Gretton et al., 2012])

Problem

- ▶ Let $S_X = \{x_i\}_{i=1}^{n_x} \sim P$ and $S_Y = \{y_i\}_{i=1}^{n_y} \sim Q$
- ▶ Question: given S_X and S_Y , is it possible to determine whether $P = Q$?
- ▶ Or, is it possible to provide a statistical test to decide $P = Q$:
 - $H_0 : P = Q$ (null hypothesis)
 - $H_A : P \neq Q$ (alternative hypothesis)

Indirect solution

1. Estimate by \hat{P} and \hat{Q} the distributions P and Q (Parzen windows, mixture of Gaussians...)
2. Compute the distance d between \hat{P} and \hat{Q} (L_2 , L_1 , KL , ...)
3. If d if beyond some threshold then reject H_0

Cons

- ▶ Curse of dimensionality
- ▶ Choice of the models to estimate the distributions
- ▶ Biased estimation of the distributions
- ▶ ...

Two-Sample Problem (see [Gretton et al., 2007, Gretton et al., 2012])


Problem

- ▶ Let $S_X = \{x_i\}_{i=1}^{n_x} \sim P$ and $S_Y = \{y_i\}_{i=1}^{n_y} \sim Q$
- ▶ Question: given S_X and S_Y , is it possible to determine whether $P = Q$?
- ▶ Or, is it possible to provide a statistical test to decide $P = Q$:
 - $H_0 : P = Q$ (null hypothesis)
 - $H_A : P \neq Q$ (alternative hypothesis)

Maximum Mean Discrepancy – MMD [Fortet and Mourier, 1953]

Find a smooth (witness) function that distinguishes P vs. Q

$$MMD(P, Q, \mathcal{F}) := \sup_{f \in \mathcal{F}} [\mathbb{E}_{X \sim P} f(X) - \mathbb{E}_{Y \sim Q} f(Y)]$$



images from D. Sejdinovic, 2014

Two-Sample Problem (see [Gretton et al., 2007, Gretton et al., 2012])



Problem

- ▶ Let $S_X = \{x_i\}_{i=1}^{n_x} \sim P$ and $S_Y = \{y_i\}_{i=1}^{n_y} \sim Q$
- ▶ Question: given S_X and S_Y , is it possible to determine whether $P = Q$?
- ▶ Or, is it possible to provide a statistical test to decide $P = Q$:
 - $H_0 : P = Q$ (null hypothesis)
 - $H_A : P \neq Q$ (alternative hypothesis)

Maximum Mean Discrepancy – MMD [Fortet and Mourier, 1953]

Find a smooth (witness) function that distinguishes P vs. Q

$$MMD(P, Q, \mathcal{F}) := \sup_{f \in \mathcal{F}} [\mathbb{E}_{X \sim P} f(X) - \mathbb{E}_{Y \sim Q} f(Y)]$$



images from D. Sejdinovic, 2014

Theoretical properties of MMD

$$MMD(P, Q, \mathcal{F}) := \sup_{f \in \mathcal{F}} [\mathbb{E}_{X \sim P} f(X) - \mathbb{E}_{Y \sim Q} f(Y)]$$

Theorem

Under mild conditions \mathcal{H} on \mathcal{F} :

$$MMD(P, Q, \mathcal{F}) = 0 \Leftrightarrow P = Q$$

Theoretical properties of MMD

$$MMD(P, Q, \mathcal{F}) := \sup_{f \in \mathcal{F}} [\mathbb{E}_{X \sim P} f(X) - \mathbb{E}_{Y \sim Q} f(Y)]$$

Theorem

Under mild conditions \mathcal{H} on \mathcal{F} :

$$MMD(P, Q, \mathcal{F}) = 0 \Leftrightarrow P = Q$$

Theorem (Dudley, 84)

\mathcal{H} : \mathcal{F} is the space of continuous bounded functions on \mathcal{X} .

Theoretical properties of MMD

$$MMD(P, Q, \mathcal{F}) := \sup_{f \in \mathcal{F}} [\mathbb{E}_{X \sim P} f(X) - \mathbb{E}_{Y \sim Q} f(Y)]$$

Theorem

Under mild conditions \mathcal{H} on \mathcal{F} :

$$MMD(P, Q, \mathcal{F}) = 0 \Leftrightarrow P = Q$$

Theorem (Dudley, 84)

\mathcal{H} : \mathcal{F} is the space of continuous bounded functions on \mathcal{X} .

Theorem (Steinwart 2001, Smola 2006)

\mathcal{H} : $\mathcal{F} = \{f : \|f\|_{\mathbb{H}} \leq 1\}$ is a unit ball of a Reproducing Kernel Hilbert Space \mathbb{H} , associated with a universal kernel. (See, e.g. Gaussian kernel, Laplace Kernel, Matern kernel.)

MMD and kernels

RKHS \mathbb{H} , kernel k

- ▶ $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ a positive kernel
- ▶ \mathbb{H} : associated RKHS

Therefore, $\forall f \in \mathbb{H}$

$$f(x) = \langle k(x, \cdot), f \rangle, \quad \forall x \in \mathcal{X}$$

(Empirical) Expectations and Kernel Mean Embedding

$$\mathbb{E}_{X \sim P} f(X) = \mathbb{E}_{X \sim P} \langle k(X, \cdot), f \rangle = \langle \mathbb{E}_{X \sim P} k(X, \cdot), f \rangle$$

$$\frac{1}{m} \sum_{i=1}^m f(X_i) = \frac{1}{m} \sum_{i=1}^m \langle k(X_i, \cdot), f \rangle = \left\langle \frac{1}{m} \sum_{i=1}^m k(X_i, \cdot), f \right\rangle$$

- ▶ $\mu(P) = \mathbb{E}_{X \sim P} k(X, \cdot) \in \mathbb{H}$ is the mean embedding of P
- ▶ $\hat{\mu}(S_X) = \frac{1}{m} \sum_{i=1}^m k(X_i, \cdot) \in \mathbb{H}$ is the empirical mean embedding of S_X

MMD and kernels

Setting $\mathcal{F} = \{f : \|f\|_{\mathbb{H}} \leq 1, f \in \mathbb{H}\}$

$$\begin{aligned} MMD(P, Q, \mathcal{F}) &= \sup_{f \in \mathcal{F}} [\mathbb{E}_{X \sim P} f(X) - \mathbb{E}_{Y \sim Q} f(Y)] \\ &= \sup_{f \in \mathcal{F}} [\langle \mathbb{E}_X k(X, \cdot), f \rangle - \langle \mathbb{E}_Y k(Y, \cdot), f \rangle] \\ &= \sup_{f \in \mathcal{F}} [\langle \mu(P), f \rangle - \langle \mu(Q), f \rangle] \\ &= \sup_{f \in \mathcal{F}} \langle \mu(P) - \mu(Q), f \rangle = \|\mu(P) - \mu(Q)\|_{\mathbb{H}} \end{aligned}$$

MMD: distance between probabilities in the kernel embedding space




image from D. Sejdinovic, 2014

MMD and kernels

In practice. . .

- ▶ $MMD^2(P, Q, \mathcal{F}) = \mathbb{E}_{X, X'} k(X, X') + \mathbb{E}_{Y, Y'} k(Y, Y') - 2\mathbb{E}_{X, Y} k(X, Y)$
- ▶ $\widehat{MMD}^2(S_X, S_Y, \mathcal{F})$ defined as

$$\frac{1}{n_x(n_x - 1)} \sum_{i \neq j} k(x_i, x_j) + \frac{1}{n_y(n_y - 1)} \sum_{i \neq j} k(y_i, y_j) - \frac{2}{n_x n_y} \sum_{i \neq j} k(x_i, y_j)$$

is an unbiased estimator of MMD^2 .

- ▶ If $n_x = n_y$, then \widehat{MMD}^2 is a U-statistic [Hoeffding, 1963, Peel et al., 2010] with “kernel” q :

$$q((x_i, y_i), (x_j, y_j)) = k(x_i, x_j) + k(y_i, y_j) - k(x_i, y_j) - k(x_j, y_i)$$

Kernel Test for the Two-Sample Problem

Observe samples $S_X = \{x_i\}_{i=1}^{n_x} \sim P$ and $S_Y = \{y_i\}_{i=1}^{n_y} \sim Q$

- ▶ $H_0 : P = Q$ (null hypothesis), $H_A : P \neq Q$ (alternative hypothesis)
- ▶ Compute the value of the statistic $\widehat{MMD}(S_X, S_Y, \mathcal{F})$ and if
 - ▶ beyond some threshold θ : reject H_0
 - ▶ otherwise: do not reject H_0

Kernel embedding wrap-up and more

One thing to remember




image from D. Sejdinovic, 2014

Extensions, uses

- ▶ Testing for independence (HSIC) [Gretton et al., 2005]
- ▶ Testing for conditional independence [Song et al., 2009]
- ▶ Testing for multivariate interaction [Sejdinovic et al., 2013]
- ▶ Kernel Bayes' rule [Fukumizu et al., 2013]
- ▶ Importance Sampling and sample bias correction [Huang et al., 2007]
- ▶ ...

Outline

Kernel Embedding of Distributions

Example: the 2-sample test problem

Kernel Embedding of Distributions

Large Scale Kernel Methods

Random Features for Kernel Machines

Fast optimization methods

Random Features for Kernel Machines

based on [[Rahimi and Recht, 2007](#), [Le et al., 2013](#)]

Kernels methods for large-scale problems

Example: Image Net Challenge

- ▶ Number of training examples $m = 1,200,000$ in ILSVRC1000 dataset.
- ▶ Dimension d of encoded features for most algorithms is more than 20,000
- ▶ Big data: d is large, m is huge

Features of Kernel Methods

	Linear kernel	Nonlinear kernel
Training speed	very fast	very slow
Training scalability	very high	low
Testing speed	very fast	very slow
Testing accuracy	lower	higher

How to get bold (using properties of the kernels considered)?

Naïve Kernel Methods

Base computational demands

- ▶ d : input feature dimension (e.g., $d = 20,000$)
- ▶ m : number of examples (e.g., $d = 1,200,000$)
- ▶ n : number of nonlinear basis function (e.g., $n = 120,000$)

Kernel expansion:

$$f(x) = \langle w, \Phi(x) \rangle = \left\langle \sum_{i=1}^n \alpha_i \Phi(x_i), \Phi(x) \right\rangle = \sum_{i=1}^n \alpha_i k(x_i, x)$$

number of features n is (often) linear with the number of examples m

	CPU Training	CPU Test	RAM Training	RAM Test
Naive	$O(m^2 d)$	$O(md)$	$O(md)$	$O(md)$

Kernel Methods with Random Features

- ▶ Random features & Linear in the feature space

- ▶ $f(x) = \sum_{i=1}^n \alpha_i \hat{\Phi}(x, w_i)$

- ▶ n random features – $n \ll m$

- ▶ $\hat{\Phi}(x, w_i)$ is a feature function obtained via randomization

	CPU Training	CPU Test	RAM Training	RAM Test
Naive	$O(m^2d)$	$O(md)$	$O(md)$	$O(md)$
Reduced set	$O(m^2d)$	$O(nd)$	$O(md)$	$O(nd)$
Low rank	$O(mnd)$	$O(nd)$	$O(nd)$	$O(nd)$
RKS	$O(mnd)$	$O(nd)$	$O(nd)$	$O(nd)$
Fastfood	$O(mn \log d)$	$O(n \log d)$	$O(n)$	$O(n)$

Random Kitchen Sinks (RKS): Approximating Kernels

Bochner's Theorem

A kernel $k(x - y)$ on \mathbb{R}^d is PSD if and only if $k(x - y)$ is the Fourier transform of a non-negative measure $p(w)$:

$$k(x - y) = \int_{\mathbb{R}^d} p(w) e^{j\langle w, (x-y) \rangle} dw$$

Consequence: approximating kernel computations

- ▶ For w_1, \dots, w_n IID according to $p(w)$

$$k(x - y) \approx \frac{1}{n} \sum_{i=1}^n e^{j\langle w_i, (x-y) \rangle} = \frac{1}{n} \sum_{i=1}^n e^{j\langle w_i, x \rangle} e^{-j\langle w_i, y \rangle} = \frac{1}{n} \sum_{i=1}^n z_{w_i}(x) z_{w_i}^*(y)$$

where $z_w(x) = \exp(j\langle w, x \rangle)$.

- ▶ If $\phi(x) = [z_{w_1}(x) \cdots z_{w_n}(x)]^\top$ then

$$k(x - y) \approx \langle \phi(x), \phi(y) \rangle.$$

- ▶ For Gaussian RBFs, we draw w from a Gaussian distribution
- ▶ Computation: $O(nd)$ CPU & $O(nd)$ RAM

Random Kitchen Sinks (RKS) with Kernel Ridge Regression

```
% Training  
W = randn(n, size(X,1));  
phi= exp(i*W*x);  
alpha = (phi*phi + eye(T)*noise)\(phi*y);  
  
% Testing  
ytest = alpha*exp(i*W*xtest);
```

from <http://www.keysduplicated.com/~ali/random-features/>

	CPU Training	CPU Test	RAM Training	RAM Test
RKS	$O(mnd)$	$O(nd)$	$O(nd)$	$O(nd)$

Massaging Wx

- ▶ Is there a way to speed up the computation of Wx ?
- ▶ Is there a way to lower the memory needs to store W ?

Answers: yes, Fastfood [Le et al., 2013]

Fastfood

Spoilers

- ▶ Wx can be computed in $O(n \log d)$ (instead of $O(nd)$)
- ▶ W need not be stored

	CPU Training	CPU Test	RAM Training	RAM Test
RKS	$O(mnd)$	$O(nd)$	$O(nd)$	$O(nd)$
Fastfood	$O(mn \log d)$	$O(n \log d)$	$O(n)$	$O(n)$

Fastfood

The magic recipe

Instead of W , use $\tilde{W} = SHG\PiHB$ where

- ▶ S is random diagonal scaling matrix (deals with spectrum)
- ▶ H is Walsh-Hadamard matrix admitting $O(d \log d)$ multiply
- ▶ G is random diagonal Gaussian matrix
- ▶ Π is random permutation matrix
- ▶ B is random binary $\{-1, 1\}$ diagonal matrix

so that \tilde{W} is approximately a Gaussian matrix

Crux: H can be built recursively (think of FFT)

$$H_{2n} = \begin{bmatrix} H_n & H_n \\ H_n & -H_n \end{bmatrix} \text{ and } H_1 = 1$$

Also

- ▶ $HG\PiHB$ produces pseudo-random Gaussian vectors (of identical length)
- ▶ S fixes the lengths to have the correct distribution

Fastfood Results

Computing $\tilde{W}x$

d	n	Fastfood	RKS	Speedup	RAM
1024	16384	0.00058s	0.0139s	24x	256x
4096	32768	0.00136s	0.1224s	90x	1024x
8196	65536	0.00268s	0.5360s	200x	2048x

Random features: partial conclusion and open questions

Conclusion

- ▶ It is possible to use features of the kernel at hand (cf. Bochner's theorem)
- ▶ Random features work well in practice
- ▶ Fastfood minimizes the storage and the computational demands

Open questions

- ▶ RKS & Fastfood transformations for other kinds of kernels?
- ▶ Things related to wavelet transforms?
- ▶ Improved generalization by implicit regularization?
- ▶ ...

Fast optimization methods

Optimization methods

- ▶ Chunking, decomposition methods
[Osuna et al., 1997, Joachims, 1998, Platt, 1998]
- ▶ Bundle methods, cutting plane methods
[Franc and Sonnenburg, 2009, Joachims et al., 2009, Teo et al., 2010]
- ▶ Active set strategies [Scheinberg, 2006]
- ▶ Stochastic gradient descent [Roux et al., 2012, Hardt et al., 2015]
- ▶ ...

Conceptual questions

- ▶ Tradeoffs of large-scale learning: optimization + approximation + estimation [Bottou and Bousquet, 2008]
- ▶ Distributed optimization and consequences on generalization
- ▶ Amortized-time integer programming
- ▶ Limited-bit representation
- ▶ ...

References I

-  Bottou, L. and Bousquet, O. (2008).
The tradeoffs of large scale learning.
In Platt, J., Koller, D., Singer, Y., and Roweis, S., editors, *Advances in Neural Information Processing Systems*, volume 20, pages 161–168. NIPS Foundation (<http://books.nips.cc>).
-  Fortet, R. and Mourier, E. (1953).
Convergence de la réparation empirique vers la réparation théorique.
Ann. Scient. École Normale Supérieure, 70:266–285.
-  Franc, V. and Sonnenburg, S. (2009).
Optimized cutting plane algorithm for large-scale risk minimization.
Journal of Machine Learning Research, 10:2157–2192.
-  Fukumizu, K., Song, L., and Gretton, A. (2013).
Kernel bayes' rule: Bayesian inference with positive definite kernels.
Journal of Machine Learning Research, 14:3753–3783.
-  Gretton, A., Borgwardt, K., Rasch, M., Schölkopf, B., and Smola, A. (2007).
A kernel method for the two-sample problem.
In *Adv in Neural Information Processing Systems 19*, pages 513–520. MIT Press.

References II

-  Gretton, A., Borgwardt, K. M., Rasch, M. J., Schlkopf, B., and Smola, A. (2012).
A kernel two-sample test.
Journal of Machine Learning Research, 13:723–773.
-  Gretton, A., Bousquet, O., Smola, A., and Schölkopf, B. (2005).
Measuring statistical dependence with hilbert-schmidt norms.
In *Proc. of Algorithmic Learning Theory*, pages 63–77. Springer-Verlag.
-  Hardt, M., Recht, B., and Singer, Y. (2015).
Train faster, generalize better: Stability of stochastic gradient descent.
-  Hoeffding, W. (1963).
Probability inequalities for sums of bounded random variables.
Journal of the American Statistical Association, 58(301):13–30.
-  Huang, J., Smola, A., Gretton, A., Borgwardt, K., and Schölkopf, B. (2007).
Correcting sample selection bias by unlabeled data.
In *Adv. in Neural Information Processing Systems 19*.

References III

-  Joachims, T. (1998).
Making Large-Scale Support Vector Machine Learning Practical.
In Schölkopf, B., Burges, C., and Smola, A., editors, *Adv. in Kernel Methods – Support Vector Learning*, pages 169–184. MIT Press, Cambridge, MA.
-  Joachims, T., Finley, T., and Yu, C.-N. (2009).
Cutting-plane training of structural svms.
Machine Learning, 77(1):27–59.
-  Le, Q., Sarlos, T., and Smola, A. (2013).
Fastfood – approximating kernel expansions in loglinear time.
In *Int. Conf. in Machine Learning*.
-  Osuna, E., Freund, R., and Girosi, F. (1997).
Improved Training Algorithm for Support Vector Machines.
In *Proc. IEEE Workshop on Neural Networks for Signal Processing*, pages 276–285.
-  Peel, T., Anthoine, S., and Ralaivola, L. (2010).
Empirical bernstein inequalities for u-statistics.
In *Advances in Neural Information Processing Systems 23 (NIPS 2010)*.

References IV

-  Platt, J. (1998).
Sequential Minimal Optimization: a Fast Algorithm for Training Support Vector Machines.
Technical Report 98-14, Microsoft Research.
-  Rahimi, A. and Recht, B. (2007).
Random features for large-scale kernel machines.
In *Adv. in Neural Information Processing Systems 19*.
-  Roux, N. L., Schmidt, M., and Bach, F. R. (2012).
A stochastic gradient method with an exponential convergence rate for finite training sets.
In *Advances in Neural Information Processing Systems 25*, pages 2663–2671.
Curran Associates, Inc.
-  Scheinberg, K. (2006).
An efficient implementation of an active set method for svm.
J. of Machine Learning Research, 7:2237–2257.
-  Sejdinovic, D., Gretton, A., and Bergsma, W. (2013).
A kernel test for three-variable interactions.
In *Adv. in Neural Information Processing Systems 26*.

References V



Song, L., Huang, J., Smola, A., and Fukumizu, K. (2009).

Hilbert space embeddings of conditional distributions with applications to dynamical systems.

In *Proc. of the 26th Int. Conf. on Machine Learning*, pages 961–968.



Teo, C. H., Vishwanathan, S. V. N., Smola, A. J., and Le, Q. V. (2010).

Bundle methods for regularized risk minimization.

Journal of Machine Learning Research, 11:311–365.

Kernel Approximation

http://www.stat.ucdavis.edu/~chohsieh/teaching/ECS289G_Fall2016/lecture9.pdf

Kernel Approximation

- Kernel methods are hard to scale up because
Kernel matrix G is an n -by- n matrix
- Can we approximate the kernel using a low-rank representation?
- Two main algorithms:
 - Nystrom approximation: Approximate the **kernel matrix**
 - Random Features: Approximate the **kernel function**

Kernel Matrix Approximation

- We want to form a low rank approximation of $G \in \mathbb{R}^{n \times n}$,
where $G_{ij} = K(\mathbf{x}_i, \mathbf{x}_j)$
- Can we do SVD?

Kernel Matrix Approximation

- We want to form a low rank approximation of $G \in \mathbb{R}^{n \times n}$,
where $G_{ij} = K(\mathbf{x}_i, \mathbf{x}_j)$
- Can we do SVD?
No, SVD needs to form the n -by- n matrix
 $O(n^2)$ space and $O(n^3)$ time

Kernel Matrix Approximation

- We want to form a low rank approximation of $G \in \mathbb{R}^{n \times n}$,
where $G_{ij} = K(\mathbf{x}_i, \mathbf{x}_j)$
- Can we do SVD?
*No, SVD needs to form the n -by- n matrix
 $O(n^2)$ space and $O(n^3)$ time*
- Can we do matrix completion or matrix sketching?

Kernel Matrix Approximation

- We want to form a low rank approximation of $G \in \mathbb{R}^{n \times n}$,
where $G_{ij} = K(\mathbf{x}_i, \mathbf{x}_j)$
- Can we do SVD?
*No, SVD needs to form the n -by- n matrix
 $O(n^2)$ space and $O(n^3)$ time*
- Can we do matrix completion or matrix sketching?
Main problem: need to generalize to new points

Nystrom Approximation

- Nystrom approximation:
 - Randomly sample m columns of G :

$$G = \begin{bmatrix} W & G_{12} \\ G_{21} & G_{22} \end{bmatrix}$$

W : m -by- m square matrix (observed)

G_{21} : $(n - m)$ -by- m matrix (observed)

$G_{12} = G_{21}^T$ (observed)

G_{22} : $(n - m)$ -by- $(n - m)$ matrix (not observed)

- Form a kernel approximation based on these mn elements

$$G \approx \tilde{G} = \begin{bmatrix} W \\ G_{21} \end{bmatrix} W^\dagger \begin{bmatrix} W & G_{12} \end{bmatrix}$$

W^\dagger : pseudo-inverse of W

Nystrom Approximation

- Why W^\dagger ?

Exact recover the top left m -by- m matrix

- The kernel approximation:

$$\begin{bmatrix} W & G_{12} \\ G_{21} & G_{22} \end{bmatrix} \approx \begin{bmatrix} W \\ G_{21} \end{bmatrix} W^\dagger \begin{bmatrix} W & G_{12} \end{bmatrix} = \begin{bmatrix} W & G_{12} \\ G_{21} & G_{21}W^\dagger G_{12} \end{bmatrix}$$

Nystrom Approximation

Algorithm:

- ① Sample m “landmark points”: $\mathbf{v}_1, \dots, \mathbf{v}_m$
- ② Compute the kernel values between all training data to landmark points:

$$C = \begin{bmatrix} W \\ G_{21} \end{bmatrix} = \begin{bmatrix} K(\mathbf{x}_1, \mathbf{v}_1) & \cdots & K(\mathbf{x}_1, \mathbf{v}_m) \\ K(\mathbf{x}_2, \mathbf{v}_1) & \cdots & K(\mathbf{x}_2, \mathbf{v}_m) \\ \vdots & \vdots & \vdots \\ K(\mathbf{x}_n, \mathbf{v}_1) & \cdots & K(\mathbf{x}_n, \mathbf{v}_m) \end{bmatrix}$$

- ③ Form the kernel approximation $\tilde{G} = CW^\dagger C^T$
(No need to explicitly form the n -by- n matrix)
- Time complexity: **mn kernel evaluations** ($O(mnd)$ time if use classical kernels)
 - How to choose m ?
Trade-off between **accuracy** v.s **computational time and memory space**

Nystrom Approximation: Training

- Solve the dual problem using low-rank representation.
- Example: Kernel ridge regression

$$\min_{\alpha} \alpha^T \tilde{G} \alpha + \lambda \|\alpha\|^2 - \mathbf{y}^T \alpha$$

Solve a linear system $(\tilde{G} + \lambda I)\alpha = \mathbf{y}$

- Use iterative methods (such as CG), with fast matrix-vector multiplication

$$(\tilde{G} + \lambda I)\mathbf{p} = CW^\dagger C^T \mathbf{p} + \lambda \mathbf{p}$$

$O(nm)$ time complexity per iteration

Nystrom Approximation: Training

- Another approach: reduce to linear ERM problems!
- Rewrite as

$$\tilde{G} = CW^\dagger C^T = C(W^\dagger)^{1/2}(W^\dagger)^{1/2}C^T$$

- So the approximate kernel can be represented by linear kernel with feature matrix $C(W^\dagger)^{1/2}$:

$$\tilde{K}(\mathbf{x}_i, \mathbf{x}_j) = \tilde{G}_{ij} = \mathbf{u}_i^T \mathbf{u}_j,$$

$$\text{where } \mathbf{u}_j = (W^\dagger)^{1/2} \begin{bmatrix} K(\mathbf{x}_j, \mathbf{v}_1) \\ \vdots \\ K(\mathbf{x}_j, \mathbf{v}_m) \end{bmatrix}$$

- The problem is equivalent to a linear models with features $\mathbf{u}_1, \dots, \mathbf{u}_n \in \mathbb{R}^m$
 - Kernel SVM \Rightarrow Linear SVM with m features
 - Kernel Ridge Regression \Rightarrow Linear Ridge regression with m features

Nystrom Approximation: Prediction

- Given the dual solution α .
- Prediction for testing sample \mathbf{x} :

$$\sum_{i=1}^n \alpha_i \tilde{K}(\mathbf{x}_i, \mathbf{x}) \quad \text{or} \quad \sum_{i=1}^n \alpha_i K(\mathbf{x}_i, \mathbf{x}) ?$$

Nystrom Approximation: Prediction

- Given the dual solution α .
- Prediction for testing sample x :

$$\sum_{i=1}^n \alpha_i \tilde{K}(x_i, x)$$

- **Need to use approximate kernel instead of original kernel!**
 - Approximate kernel gives much better accuracy!
 - Because training & testing should use the same kernel.
- Generalization bound:
 - (Alaoui and Mahoney, "Fast Randomized Kernel Methods With Statistical Guarantees". 2014.)
 - (Rudi et al., "Less is More: Nystrom Computational Regularization". NIPS 2015.)
 - (using small number of landmark points can be viewed as a regularization)

Nystrom Approximation: Prediction

- How to define $\tilde{K}(\mathbf{x}, \mathbf{x}_i)$?

$$\tilde{G} = \begin{bmatrix} K(\mathbf{x}_1, \mathbf{v}_1) & \cdots & K(\mathbf{x}_1, \mathbf{v}_m) \\ \vdots & \vdots & \vdots \\ K(\mathbf{x}_n, \mathbf{v}_1) & \cdots & K(\mathbf{x}_n, \mathbf{v}_m) \\ \mathbf{K}(\mathbf{x}, \mathbf{v}_1) & \cdots & \mathbf{K}(\mathbf{x}, \mathbf{v}_m) \end{bmatrix} W^\dagger \begin{bmatrix} K(\mathbf{v}_1, \mathbf{x}_1) & \cdots & K(\mathbf{v}_1, \mathbf{x}) \\ \vdots & \vdots & \vdots \\ K(\mathbf{v}_m, \mathbf{x}_1) & \cdots & K(\mathbf{v}_m, \mathbf{x}) \end{bmatrix}$$

- Compute $\mathbf{u} = (W^\dagger)^{1/2} \begin{bmatrix} K(\mathbf{v}_1, \mathbf{x}) \\ \vdots \\ K(\mathbf{v}_m, \mathbf{x}) \end{bmatrix}$ (need m kernel evaluations)
- Approximate kernel can be defined by

$$\tilde{K}(\mathbf{x}, \mathbf{x}_i) = \mathbf{u}^T \mathbf{u}_i$$

Nystrom Approximation: Prediction

- Prediction:

$$f(\mathbf{x}) = \sum_{i=1}^n \alpha_i \mathbf{u}^T \mathbf{u}_i = \mathbf{u}^T \left(\sum_{i=1}^n \alpha_i \mathbf{u}_i \right)$$

$\left(\sum_{i=1}^n \alpha_i \mathbf{u}_i \right)$ can be pre-computed

⇒ prediction time is m kernel evaluations

- Original kernel method: need n kernel evaluations for prediction.
- Summary:

	Quality	Training time	Prediction time
Original kernel	exact	$n^{1.x}$ kernel + kernel training	n kernel
Nystrom (m)	rank m	nm kernel + linear training	m kernel

Nystrom Approximation

- How to select landmark points ($\mathbf{v}_1, \dots, \mathbf{v}_m$)
 - Traditional approach: uniform random sampling from training data
 - Importance sampling (leverage score):
 - Drineas and Mahoney “On the Nystrom method for approximating a Gram matrix for improved kernel-based learning”. JMLR 2015.
 - Gittens and Mahoney “Revisiting the Nystrom method for improved large-scale machine learning”. ICML 2013
 - Kmeans sampling (performs extremely well) :
Run kmeans clustering and set $\mathbf{v}_1, \dots, \mathbf{v}_m$ to be cluster centers
 - Zhang et al., “Improved Nystrom low rank approximation and error analysis”. ICML 2008.
 - Subspace distance:
 - Lim et al., “Double Nystrom method: An efficient and accurate Nystrom scheme for large-scale data sets”. ICML 2015.

Nystrom Approximation (other related papers)

- Distributed setting: (Kumar et al., “Ensemble Nystrom method”. NIPS 2009).
- Block diagonal + low-rank approximation: (Si et al., “Memory efficient kernel approximation”. ICML 2014.)
- Nystrom method for fast prediction: (Hsieh et al., “Fast prediction for large-scale kernel machines. ” NIPS, 2014.)
- Structured landmark points: (Si et al., “Computational Efficient Nystrom Approximation using Fast Transforms”. ICML 2016.)